# DONG-WON LEE

linkedin.com/in/dong-won-lee | dlehddnjs245@khu.ac.kr

## EDUCATION

**Kyung Hee University** Mar 2021 – Present, Expected Aug 2027
*B.S. in Software Convergence | GPA: 3.91 / 4.50*

## EXPERIENCE

**EMIL Lab, Kyung Hee University** Jul 2026 – Present
*Undergraduate Research Intern*
- Studying memory-efficient KV cache quantization for reasoning LLMs (in progress).

**ROBROS** Dec 2025 – Mar 2026
*Software Engineer Internship*
- Built real-time multi-camera streaming for a humanoid robot SDK over CycloneDDS, from an on-robot NVIDIA Jetson through a bridge to client PCs across a mixed LAN/Wi-Fi network (Python, partly C++).
- Set up cross-network routing (IP forwarding, NAT), tuned UDP socket buffers, and shipped it as a systemd service with a Cython-compiled core; validated on the physical robot.

**Kyung Hee University** Jul 2025 – Dec 2025
*Biometric Data Analysis Research (Research Assistant)*
- Social data analysis of ChatGPT use in elementary-school classrooms: built a multimodal classroom-interaction dataset, defined annotation criteria, checked consistency, and ran hypothesis-driven analysis (Python).

**Photo Studio Booking Platform** Mar 2024 – Jul 2025
*Student Founder & Product Manager*
- Founded a student-run startup building a photo studio and photographer search/booking app; led a 5-member team across product, design, and performance marketing.
- Built and shipped the MVP as an iOS app plus an admin web (React Native, Node.js, GCP, GitHub Actions); reached 16 B2B partner studios and photographers and 500+ DAU, and monetized the app on the App Store.

## PROJECTS

**Memory-Efficient Long-Video Understanding** Jun 2026 – Present
*Capstone Project*
- On top of InfiniPot-V KV cache compression for streaming long video (Qwen2-VL-7B), storing the kept KV cache in INT4 instead of FP16 to reduce KV memory at the same token budget.
- Evaluating on Video-MME and MLVU against FP16 baselines.

**GPU-Accelerated LLM Inference** Mar 2025 – Present
- Implemented GPU kernels with core optimization techniques to understand hardware-level bottlenecks.
- Built self-hosted LLM serving with vLLM on limited GPU resources.

**Data-Efficient 3D Reconstruction** Mar 2026 – Jun 2026
- Input-view selection for 3D Gaussian Splatting via dataset distillation, comparing view-scoring strategies against uniform/random baselines.

## MILITARY SERVICE

**Republic of Korea Army** Jul 2022 – Jan 2024
*Completed military service as Army Sergeant*

## SKILLS

**Languages:** Python, C++, CUDA

**ML / Inference:** PyTorch, vLLM, Model Quantization, LLM Serving